



# State Farm Distracted Driver Detection

Made By:

- Mohamed Khalifa
- Sara Samer
- Amira Adel
- Hadeer Mahdy
- Ahmed Gamal

# Dataset

The dataset used in this project was provided by State Farm through a Kaggle competition, which is a set of images of drivers taken inside a car capturing their activities such as texting, talking on the phone, eating, reaching behind, putting on makeup, etc. As shown in Figure 1. These activities are classified into 10 classes as:

- c0: safe driving
- c1: texting — right
- c2: talking on the phone — right
- c3: texting — left
- c4: talking on the phone — left
- c5: operating the radio
- c6: drinking
- c7: reaching behind
- c8: hair and makeup
- c9: talking to a passenger

# Exploratory Visualization

talking on the phone - left



reaching behind



safe driving



texting - left



texting - right



texting - left



texting - left



talking on the phone - right



talking on the phone - right

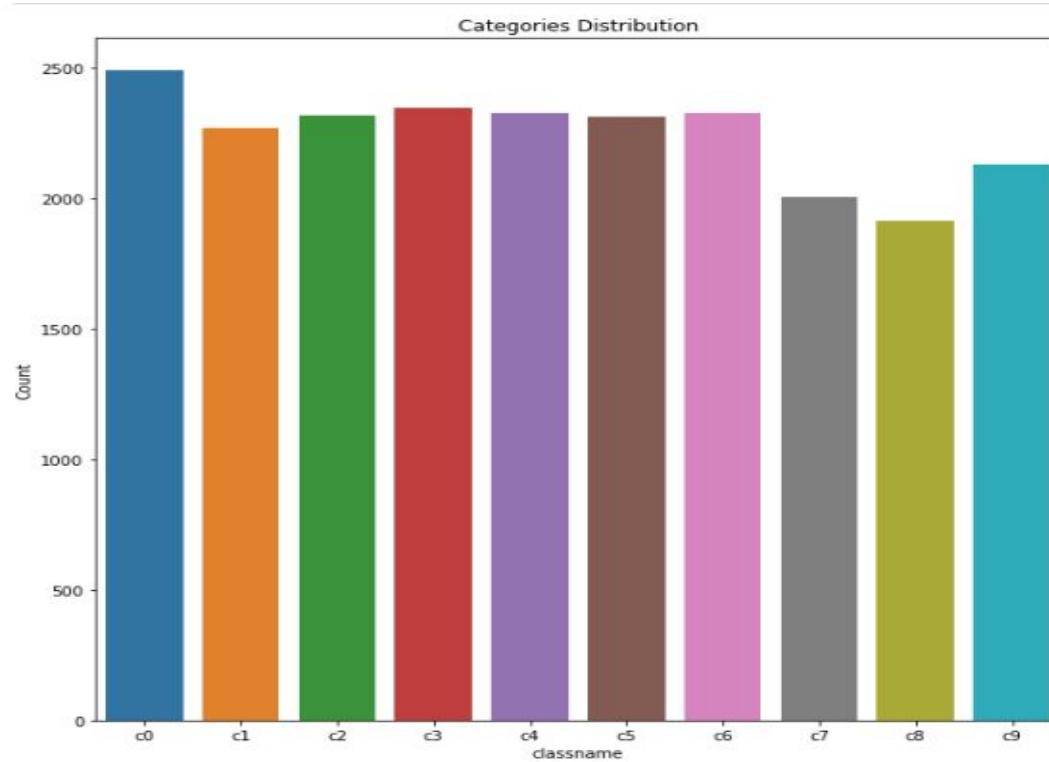


drinking



# Data Exploration

There is one characteristic about the input data that may need to be addressed. I noticed that some categories have more images than the others as we can see from the chart below.



# 1. Data Preprocessing

These are the data preprocessing steps that positively impacted results and were applied on the data:

1. Images are converted into 3 channels, RGB – This preprocessing is done so models may use the 3 channels to learn features with the objective of improving accuracy and log loss.
2. Images are resized to 224 x 224 – Resizing images makes it easier to load on memory at the cost of losing some details.
3. The list of images are normalized and divided into a train set and validation set - This division is important so that we could validate if our model is improving or not when it is training.

# Statistics

```
# Statistics
# Load the list of names
names = [item[17:19] for item in sorted(glob("../input/state-farm-distracted-driver-detection/imgs/train/*"))]
test_files_size = len(np.array(glob(os.path.join('..', 'input', '../input/state-farm-distracted-driver-detection/imgs/test/*', '*.jpg'))))
x_train_size = len(x_train)
categories_size = len(names)
x_test_size = len(x_test)
print('There are %s total images.\n' % (test_files_size + x_train_size + x_test_size))
print('There are %d training images.' % x_train_size)
print('There are %d total training categories.' % categories_size)
print('There are %d validation images.' % x_test_size)
print('There are %d test images.' % test_files_size)
```

There are 102150 total images.

There are 17939 training images.

There are 10 total training categories.

There are 4485 validation images.

There are 79726 test images.

## 2. Baseline Dense Model

Model summary:

---

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 512)	25690624
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 10)	1290
<hr/>		

Total params: 25,856,138

Trainable params: 25,856,138

Non-trainable params: 0

# Model accuracy:

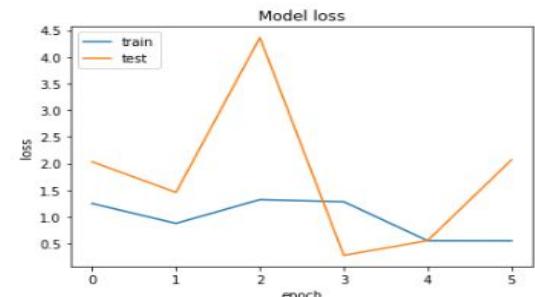
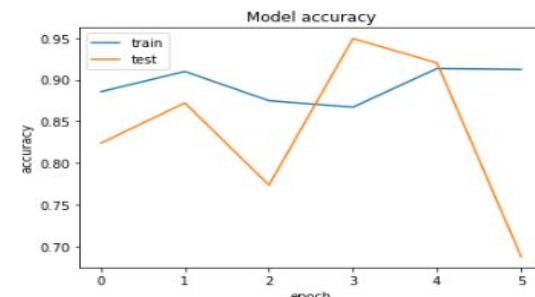
```
# Training the Dense Model version
history_v0 = model_v0.fit(x_train.reshape(17939,224*224*1), y_train,
                           validation_data=(x_test.reshape(4485,224*224*1), y_test),
                           callbacks=callbacks,
                           epochs=nb_epoch, batch_size=batch_size, verbose=1)

Epoch 1/10
449/449 [=====] - 4s 9ms/step - loss: 1.2520 - accuracy: 0.8857 - val_loss: 2.0341 - val_accuracy: 0.8239
Epoch 00001: val_loss did not improve from 1.98025
Epoch 2/10
449/449 [=====] - 3s 8ms/step - loss: 0.8795 - accuracy: 0.9099 - val_loss: 1.4608 - val_accuracy: 0.8722
Epoch 00002: val_loss improved from 1.98025 to 1.46082, saving model to /saved_models/weights_best_vanilla.hdf5
Epoch 3/10
449/449 [=====] - 3s 7ms/step - loss: 1.3258 - accuracy: 0.8749 - val_loss: 4.3618 - val_accuracy: 0.7735
Epoch 00003: val_loss did not improve from 1.46082
Epoch 4/10
449/449 [=====] - 3s 7ms/step - loss: 1.2856 - accuracy: 0.8671 - val_loss: 0.2810 - val_accuracy: 0.9494
Epoch 00004: val_loss improved from 1.46082 to 0.28098, saving model to /saved_models/weights_best_vanilla.hdf5
Epoch 5/10
449/449 [=====] - 3s 8ms/step - loss: 0.5544 - accuracy: 0.9135 - val_loss: 0.5580 - val_accuracy: 0.9202
Epoch 00005: val_loss did not improve from 0.28098
Epoch 6/10
449/449 [=====] - 3s 7ms/step - loss: 0.5534 - accuracy: 0.9124 - val_loss: 2.0716 - val_accuracy: 0.6881
Epoch 00006: val_loss did not improve from 0.28098
Epoch 00006: early stopping
```

```
score = model_v0.evaluate(x_test.reshape(4485,224*224*1), y_test, verbose=1)
print('Score: ', score)
```

```
141/141 [=====] - 0s 3ms/step - loss: 2.0716 - accuracy: 0.6881
Score: [2.0716123580932617, 0.6880713701248169]
```

```
plot_train_history(history_v0)
```



# 3. Baseline CNN

Model summary:

flatten_2 (Flatten)	(None, 8192)	0
dense_14 (Dense)	(None, 512)	4194816
batch_normalization_20 (Batch Normalization)	(None, 512)	2048
dropout_13 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 128)	65664
dropout_14 (Dropout)	(None, 128)	0
dense_16 (Dense)	(None, 10)	1290
<hr/>		
Total params:	4,552,042	
Trainable params:	4,550,122	
Non-trainable params:	1,920	

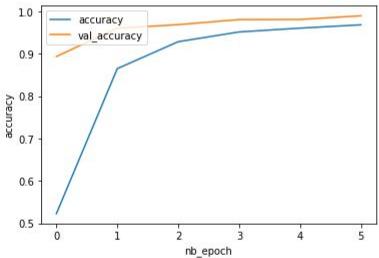
Model: "sequential_5"		
Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 62, 62, 32)	320
batch_normalization_14 (Batch Normalization)	(None, 62, 62, 32)	128
conv2d_13 (Conv2D)	(None, 62, 62, 32)	9248
batch_normalization_15 (Batch Normalization)	(None, 62, 62, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 31, 31, 32)	0
dropout_10 (Dropout)	(None, 31, 31, 32)	0
conv2d_14 (Conv2D)	(None, 31, 31, 64)	18496
batch_normalization_16 (Batch Normalization)	(None, 31, 31, 64)	256
conv2d_15 (Conv2D)	(None, 31, 31, 64)	36928
batch_normalization_17 (Batch Normalization)	(None, 31, 31, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_11 (Dropout)	(None, 16, 16, 64)	0
conv2d_16 (Conv2D)	(None, 16, 16, 128)	73856
batch_normalization_18 (Batch Normalization)	(None, 16, 16, 128)	512
conv2d_17 (Conv2D)	(None, 16, 16, 128)	147584
batch_normalization_19 (Batch Normalization)	(None, 16, 16, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 8, 8, 128)	0
dropout_12 (Dropout)	(None, 8, 8, 128)	0

# Model accuracy:

```
plt.plot(history_v2.history['accuracy'])
plt.plot(history_v2.history['val_accuracy'])

plt.ylabel('accuracy')
plt.xlabel('nb_epoch')
plt.legend(['accuracy', 'val_accuracy'], loc='upper left')

plt.show()
```



+ Code + Markdown

```
score1 = model.evaluate(x_test, y_test, verbose=1)
print('Loss: ', score1[0])
print('Accuracy: ', score1[1]*100, '%')
```

```
141/141 [=====] - 15s 103ms/step - loss: 0.0351 - accuracy: 0.9897
Loss: 0.035131312906742096
Accuracy: 98.97435903549194 %
```

```
history_v2 = model.fit(x_train, y_train,
                       validation_data=(x_test, y_test),
                       epochs=nb_epoch, batch_size=batch_size, verbose=1)
```

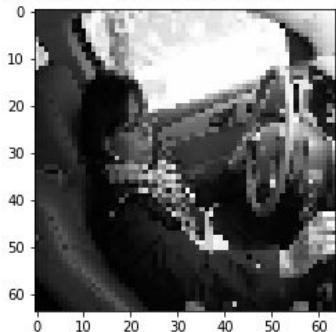
```
Epoch 1/6
281/281 [=====] - 220s 771ms/step - loss: 1.5443 - accuracy: 0.4926 - val_loss: 0.4259 - val_accuracy: 0.8658
Epoch 2/6
281/281 [=====] - 213s 757ms/step - loss: 0.4442 - accuracy: 0.8537 - val_loss: 0.1538 - val_accuracy: 0.9532
Epoch 3/6
281/281 [=====] - 209s 745ms/step - loss: 0.2455 - accuracy: 0.9216 - val_loss: 0.0702 - val_accuracy: 0.9808
Epoch 4/6
281/281 [=====] - 209s 745ms/step - loss: 0.1742 - accuracy: 0.9451 - val_loss: 0.0656 - val_accuracy: 0.9824
Epoch 5/6
281/281 [=====] - 212s 755ms/step - loss: 0.1293 - accuracy: 0.9589 - val_loss: 0.0574 - val_accuracy: 0.9815
Epoch 6/6
281/281 [=====] - 212s 755ms/step - loss: 0.0985 - accuracy: 0.9686 - val_loss: 0.0421 - val_accuracy: 0.9882
```

# Prediction Samples:

1/1 [=====] - 0s 370ms/step

Y prediction: [[1.9079283e-09 8.7641758e-16 2.4082774e-11 5.5070032e-10 2.4784116e-10  
9.9999964e-01 3.3204529e-12 8.7869495e-10 9.7563579e-08 2.1145745e-07]]

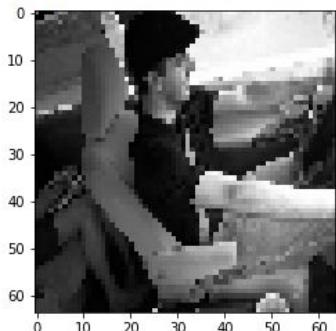
Predicted: Operating the radio



1/1 [=====] - 0s 50ms/step

Y prediction: [[9.6494182e-07 1.5600721e-09 6.6973621e-07 4.2936104e-09 2.5467530e-06  
9.9985135e-01 4.9533831e-08 5.3989105e-08 1.6404088e-06 1.4281916e-04]]

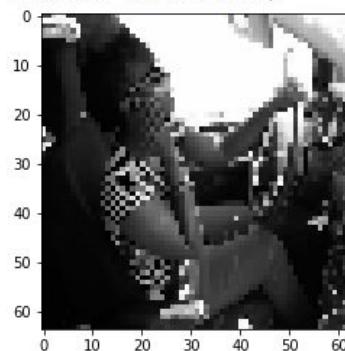
Predicted: Operating the radio



1/1 [=====] - 0s 51ms/step

Y prediction: [[1.5194789e-06 6.5872609e-06 6.5887538e-03 6.2759011e-09 2.4896980e-07  
1.8446251e-08 2.3642656e-07 6.7874826e-06 9.9319708e-01 1.9877782e-04]]

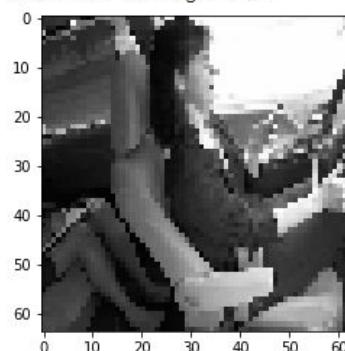
Predicted: Hair and makeup



1/1 [=====] - 0s 45ms/step

Y prediction: [[1.5350214e-03 1.5164029e-07 1.6971776e-06 5.8196956e-01 2.9925475e-01  
9.9514835e-02 1.1400979e-07 2.0523319e-06 3.2206574e-03 1.4501143e-02]]

Predicted: Texting - left



# 4. Data Augmentation

Data generator:

```
train_datagen = ImageDataGenerator(rescale = 1.0/255,  
                                  shear_range = 0.2,  
                                  zoom_range = 0.2,  
                                  horizontal_flip = True,  
                                  validation_split = 0.2)  
  
test_datagen = ImageDataGenerator(rescale=1.0/ 255, validation_split = 0.2)
```

# Model Accuracy:

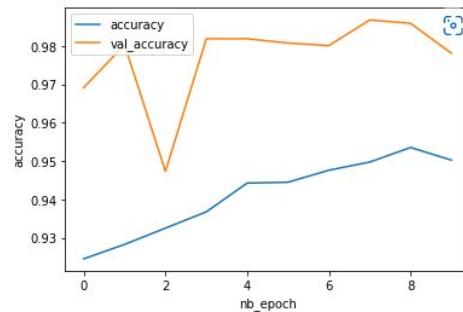
```
history_v1 = model_1.fit_generator(training_generator,
                                    steps_per_epoch = nb_train_samples // batch_size,
                                    epochs = 10,
                                    verbose = 1,
                                    callbacks=callbacks,
                                    validation_data = validation_generator,
                                    validation_steps = nb_validation_samples // batch_size)
```

```
/opt/conda/lib/python3.7/site-packages/keras/engine/training.py:1972: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

Epoch 1/10
280/280 [=====] - 219s 781ms/step - loss: 0.2360 - accuracy: 0.9245 - val_loss: 0.1119 - val_accuracy: 0.9692
Epoch 2/10
280/280 [=====] - 222s 793ms/step - loss: 0.2268 - accuracy: 0.9283 - val_loss: 0.0595 - val_accuracy: 0.9804
Epoch 3/10
280/280 [=====] - 220s 786ms/step - loss: 0.2088 - accuracy: 0.9325 - val_loss: 0.1863 - val_accuracy: 0.9473
Epoch 4/10
280/280 [=====] - 219s 781ms/step - loss: 0.1983 - accuracy: 0.9368 - val_loss: 0.0495 - val_accuracy: 0.9819
Epoch 5/10
280/280 [=====] - 218s 778ms/step - loss: 0.1776 - accuracy: 0.9443 - val_loss: 0.0545 - val_accuracy: 0.9819
Epoch 6/10
280/280 [=====] - 223s 796ms/step - loss: 0.1712 - accuracy: 0.9445 - val_loss: 0.0641 - val_accuracy: 0.9808
Epoch 7/10
280/280 [=====] - 220s 786ms/step - loss: 0.1645 - accuracy: 0.9476 - val_loss: 0.0604 - val_accuracy: 0.9801
Epoch 8/10
280/280 [=====] - 218s 780ms/step - loss: 0.1560 - accuracy: 0.9498 - val_loss: 0.0425 - val_accuracy: 0.9868
Epoch 9/10
280/280 [=====] - 219s 782ms/step - loss: 0.1563 - accuracy: 0.9536 - val_loss: 0.0476 - val_accuracy: 0.9859
Epoch 10/10
280/280 [=====] - 225s 803ms/step - loss: 0.1555 - accuracy: 0.9503 - val_loss: 0.0638 - val_accuracy: 0.9781
```

# Model Evaluation:

```
:  
    plt.plot(history_v1.history["accuracy"])  
    plt.plot(history_v1.history['val_accuracy'])  
  
    plt.ylabel('accuracy')  
    plt.xlabel('nb_epoch')  
    plt.legend(['accuracy', 'val_accuracy'], loc='upper left')  
  
    plt.show()
```



+ Code

+ Markdown

```
:  
    score2 = model_1.evaluate_generator(validation_generator, nb_validation_samples // batch_size)  
    print('Loss: ', score2[0])  
    print('Accuracy: ', score2[1]*100, '%')
```

Loss: 0.06373465061187744  
Accuracy: 97.81249761581421 %

# 5. Transfer Learning

## Model Selection:

### ▼ Model development

```
[9] base = VGG16(weights='imagenet', include_top=False, input_shape=(256, 256, 3))  
    base.trainable = False
```

# Model Summary:

Model: "vgg16"		
Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[None, 256, 256, 3]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808

block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0

Total params: 14,714,688  
Trainable params: 0  
Non-trainable params: 14,714,688

# Added Layers:

```
model.summary()

Model: "sequential"
-----

| Layer (type)       | Output Shape      | Param #  |
|--------------------|-------------------|----------|
| vgg16 (Functional) | (None, 8, 8, 512) | 14714688 |
| flatten (Flatten)  | (None, 32768)     | 0        |
| dropout (Dropout)  | (None, 32768)     | 0        |
| dense (Dense)      | (None, 256)       | 8388864  |
| dense_1 (Dense)    | (None, 128)       | 32896    |
| dense_2 (Dense)    | (None, 10)        | 1290     |


-----  
Total params: 23,137,738  
Trainable params: 8,423,050  
Non-trainable params: 14,714,688
```

# Model Accuracy:

## Adam Opt.

```
history = model.fit(  
    train_generator,  
    steps_per_epoch=560,  
    epochs=5,  
    validation_data=val_generator,  
    validation_steps=140  
)  
  
Epoch 1/5  
560/560 [=====] - 181s 312ms/step - loss: 0.3464 - acc: 0.9185 - val_loss: 0.0675 - val_acc: 0.9799  
Epoch 2/5  
560/560 [=====] - 174s 310ms/step - loss: 0.0453 - acc: 0.9863 - val_loss: 0.0465 - val_acc: 0.9875  
Epoch 3/5  
560/560 [=====] - 173s 309ms/step - loss: 0.0320 - acc: 0.9900 - val_loss: 0.0351 - val_acc: 0.9893  
Epoch 4/5  
560/560 [=====] - 171s 305ms/step - loss: 0.0372 - acc: 0.9889 - val_loss: 0.0498 - val_acc: 0.9850  
Epoch 5/5  
560/560 [=====] - 172s 307ms/step - loss: 0.0395 - acc: 0.9872 - val_loss: 0.0158 - val_acc: 0.9955
```

# Model Accuracy:

## RMSprop Opt.

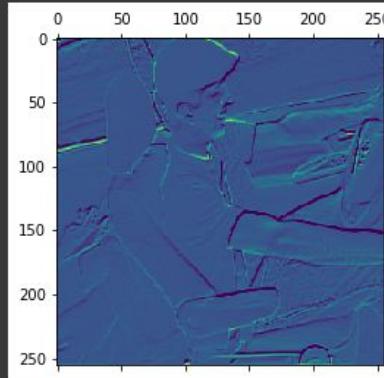
```
[ ] history = model2.fit(  
    train_generator,  
    steps_per_epoch=560,  
    epochs=7,  
    validation_data=val_generator,  
    validation_steps=140  
)  
  
Epoch 1/7  
560/560 [=====] - 173s 308ms/step - loss: 0.5192 - acc: 0.8390 - val_loss: 0.0348 - val_acc: 0.9913  
Epoch 2/7  
560/560 [=====] - 173s 310ms/step - loss: 0.1117 - acc: 0.9700 - val_loss: 0.0305 - val_acc: 0.9911  
Epoch 3/7  
560/560 [=====] - 173s 309ms/step - loss: 0.0809 - acc: 0.9812 - val_loss: 0.3306 - val_acc: 0.9228  
Epoch 4/7  
560/560 [=====] - 171s 305ms/step - loss: 0.0575 - acc: 0.9867 - val_loss: 0.0269 - val_acc: 0.9935  
Epoch 5/7  
560/560 [=====] - 175s 311ms/step - loss: 0.0630 - acc: 0.9882 - val_loss: 0.1178 - val_acc: 0.9714  
Epoch 6/7  
560/560 [=====] - 173s 308ms/step - loss: 0.0430 - acc: 0.9913 - val_loss: 0.0223 - val_acc: 0.9953  
Epoch 7/7  
560/560 [=====] - 176s 315ms/step - loss: 0.0436 - acc: 0.9927 - val_loss: 0.0691 - val_acc: 0.9862
```

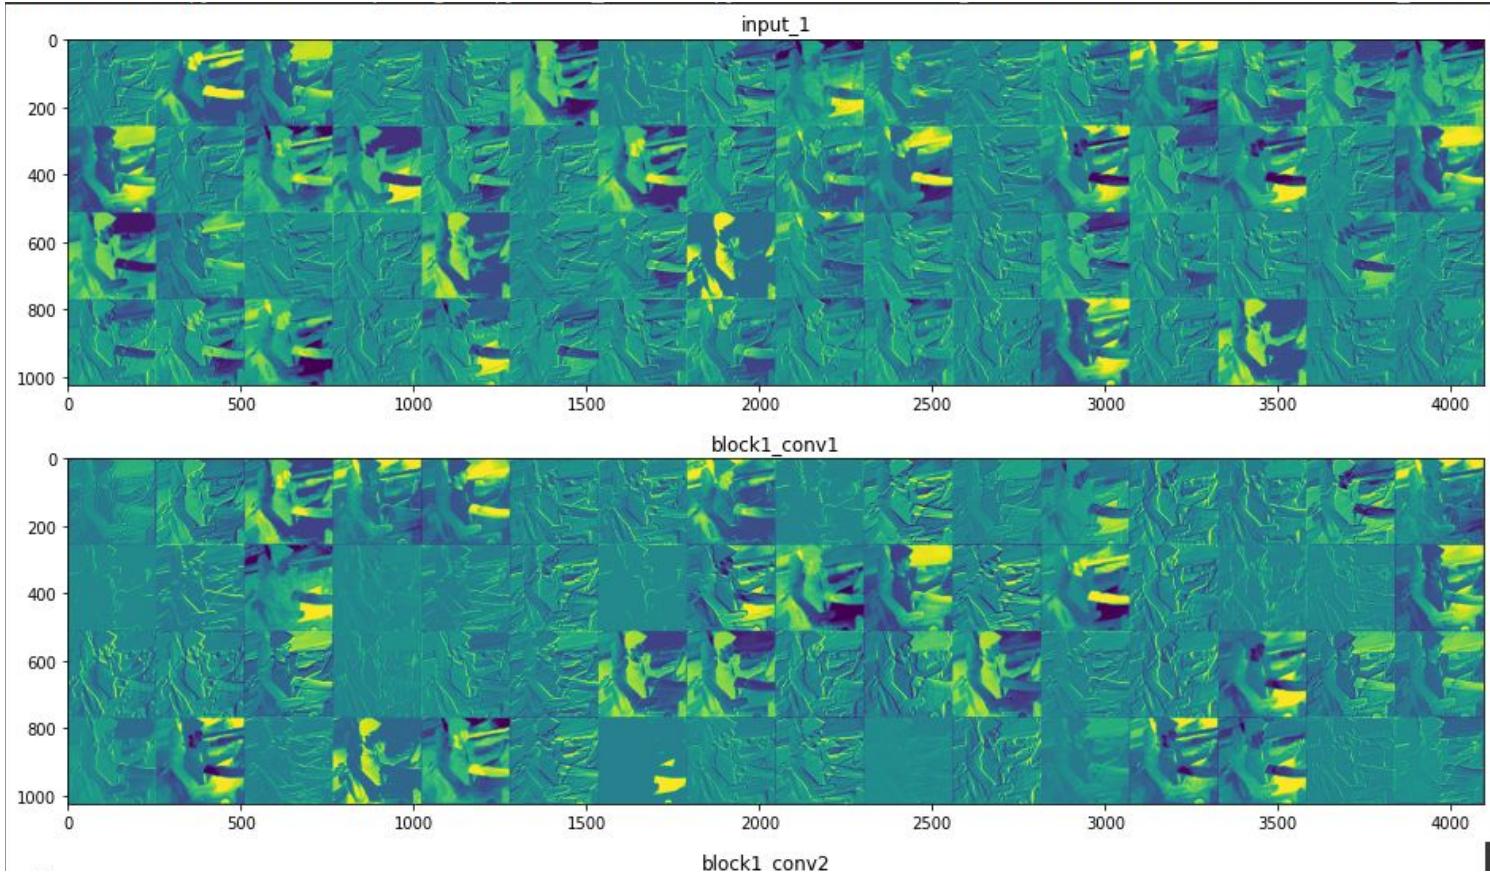
# 6. CNN Visualization

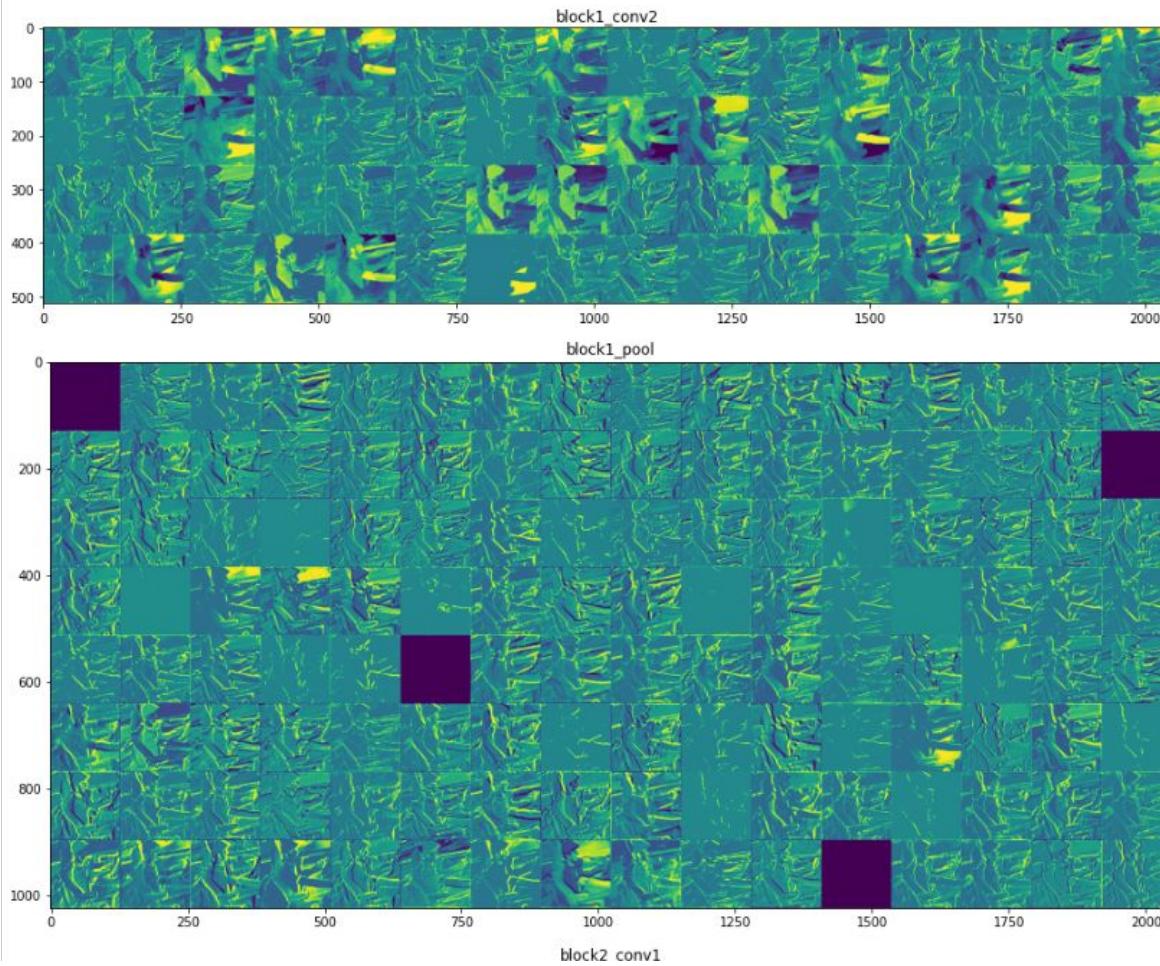
```
[ ] img = image.load_img('/content/gdrive/MyDrive/DL/test.jpg', target_size=(256, 256))
img_tensor = image.img_to_array(img)
img_tensor = np.expand_dims(img_tensor, axis=0)
img_tensor /= 255.
print(img_tensor.shape)
plt.matshow(img)
plt.show()
```



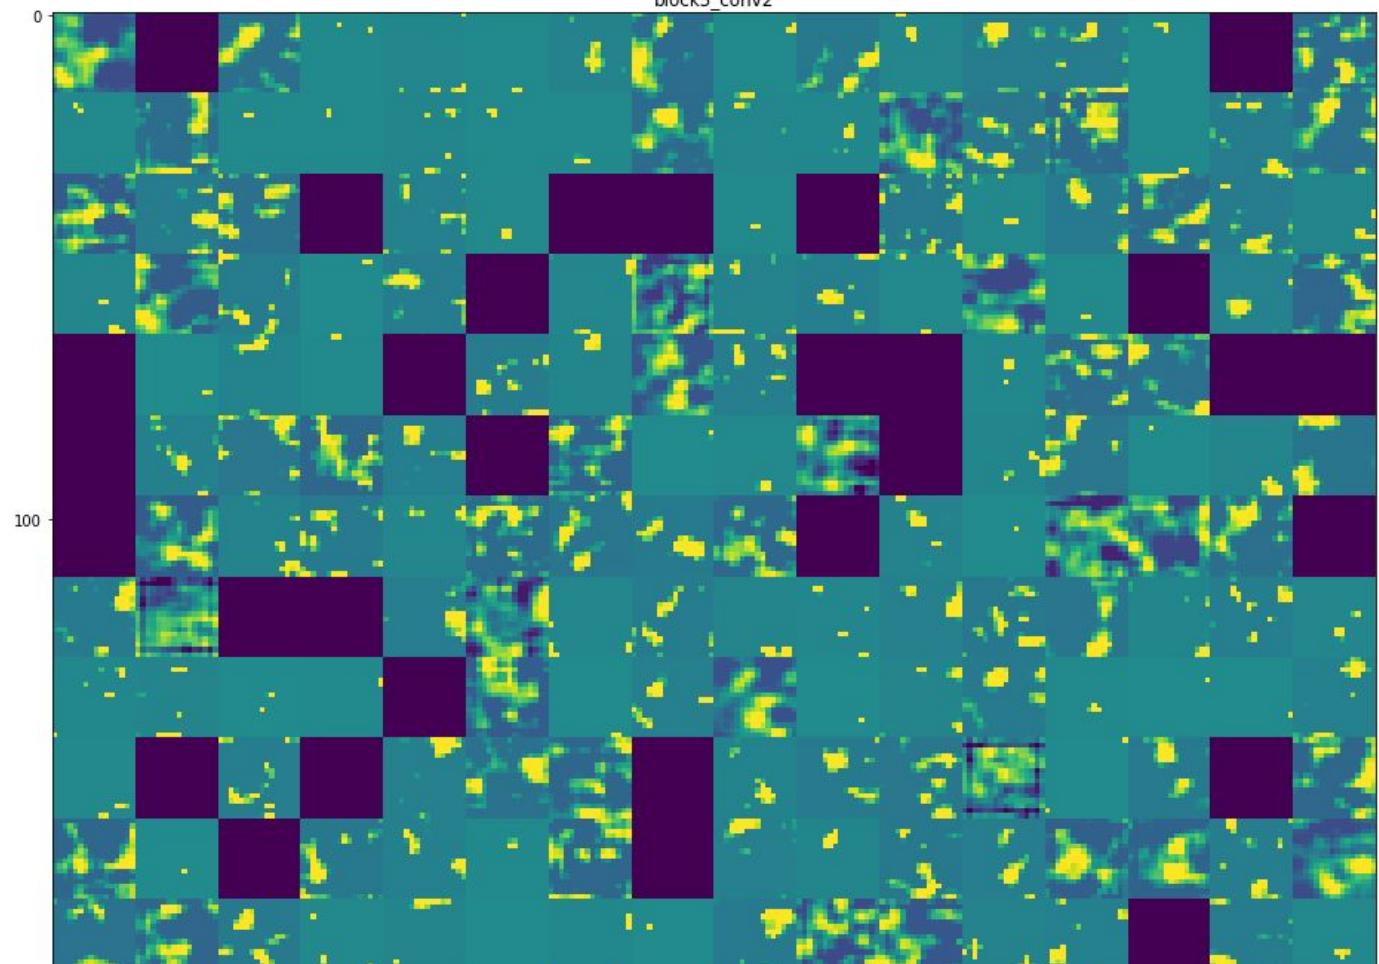
```
[ ] activations = activation_model.predict(img_tensor)
first_layer_activation = activations[0]
plt.matshow(first_layer_activation[0, :, :, 3], cmap='viridis')
plt.show()
```







block5\_conv2



# HeatMap

